

# Who I Am

- Daniel Leech
- From England
- Open Source Developer
- Bicycle Tourer

# Open Source Developer

- Symfony CMF Core Team
- Sulu CMS
- PHPCR-ODM
- CMF Routing Auto
- PHPCR Shell
- Jackalope FS



**doctrine**



# Bicycle Tour 2015

# Austria



# Slovakia





# Hungary



# Romania



# Bulgaria





# Turkey



# Greece





# Albania





# Montenegro





# Bosnia and Herzegovina



# Zagreb!



# Why Benchmarking?

Jackalope 2

*(this is not the real logo)*





# Benchmarking





# What is Benchmarking?

“In computing, a benchmark is the act of running a computer program, a set of programs, or other operations, in order to assess the ***relative*** performance of an object, normally by running a number of ***standard tests*** and trials against it”

- Wikipedia

# What I Talk About When I Talk About Benchmarking

“The measurement of time taken to execute some PHP code compared to the same measurement of equivalent PHP code”

- Me

# Microtime!

```
$start = microtime(true);  
    // do something  
$timeTaken = microtime(true) - $start;  
echo 'Time 1: ' . $timeTaken;
```

```
$start = microtime(true);  
    // do something else  
$timeTaken = microtime(true) - $start;  
echo 'Time 2: ' . $timeTaken;
```

# About Microtime

- Amount of time elapsed since the UNIX epoch accurate to a microsecond.
- The symbol for microsecond is  $\mu\text{s}$  (Greek “mu”)
- 1 microsecond = 1,000,000th ( $10^6$ ) of a second!
- Highest degree of precision available in PHP.



# Microseconds


- Human eye blink takes 350,000 microseconds (just over 1/4 of one second).
- Human finger click takes 150,000 microseconds (just over 1/7 of one second).
- Some PHP functions execute  $< 1\mu s$



# Why Would You Want to Benchmark?

- Compare
  - time?
  - memory?
- What
  - algorithms
  - hardware
  - platforms
  - dependencies
- Prevent
  - regressions
- Satisfy
  - your curiosity

WIN



ARGUMENTS

(OBJECTIVELY)

# Comparing Stuff



Compare

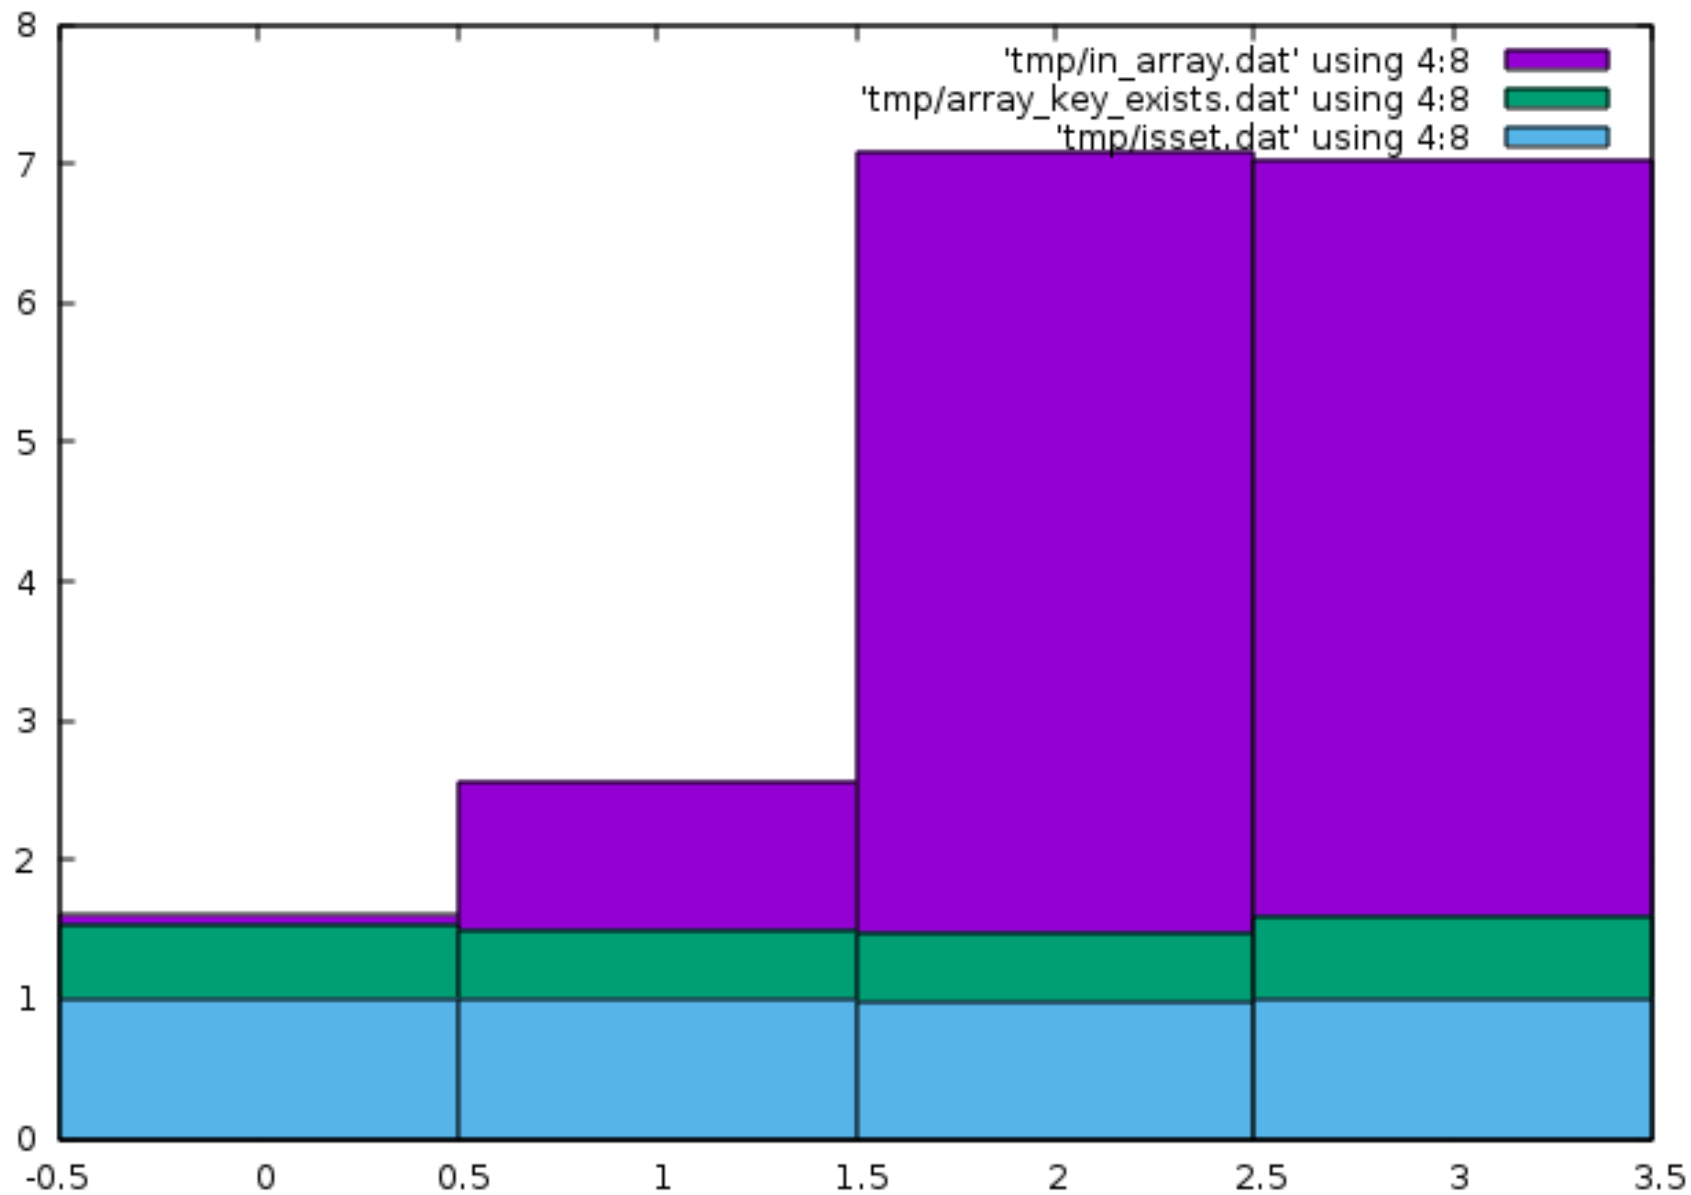
```
isset($array['foo'])
```

Functions



```
array_key_exists('foo', $array)
```





Elements

10

100

1000

10,000



*Compare*



**git** master

**VS**

*Branches*



**git** harder\_better\_stronger\_faster

**FIGHT**




Compare



MariaDB

Implementations


VS



PostgreSQL

FIGHT

# Find 100 entities with Doctrine ORM




Compare

Platforms



# md5: PHP3 vs. PHP5 vs. PHP7





Compare

Hardware

VS



FIGHT

# Empirical Knowledge





# Considerations



# Considerations

- Major: Stability
- Minor: Observer effect
- Minor: Recovery

# Stability

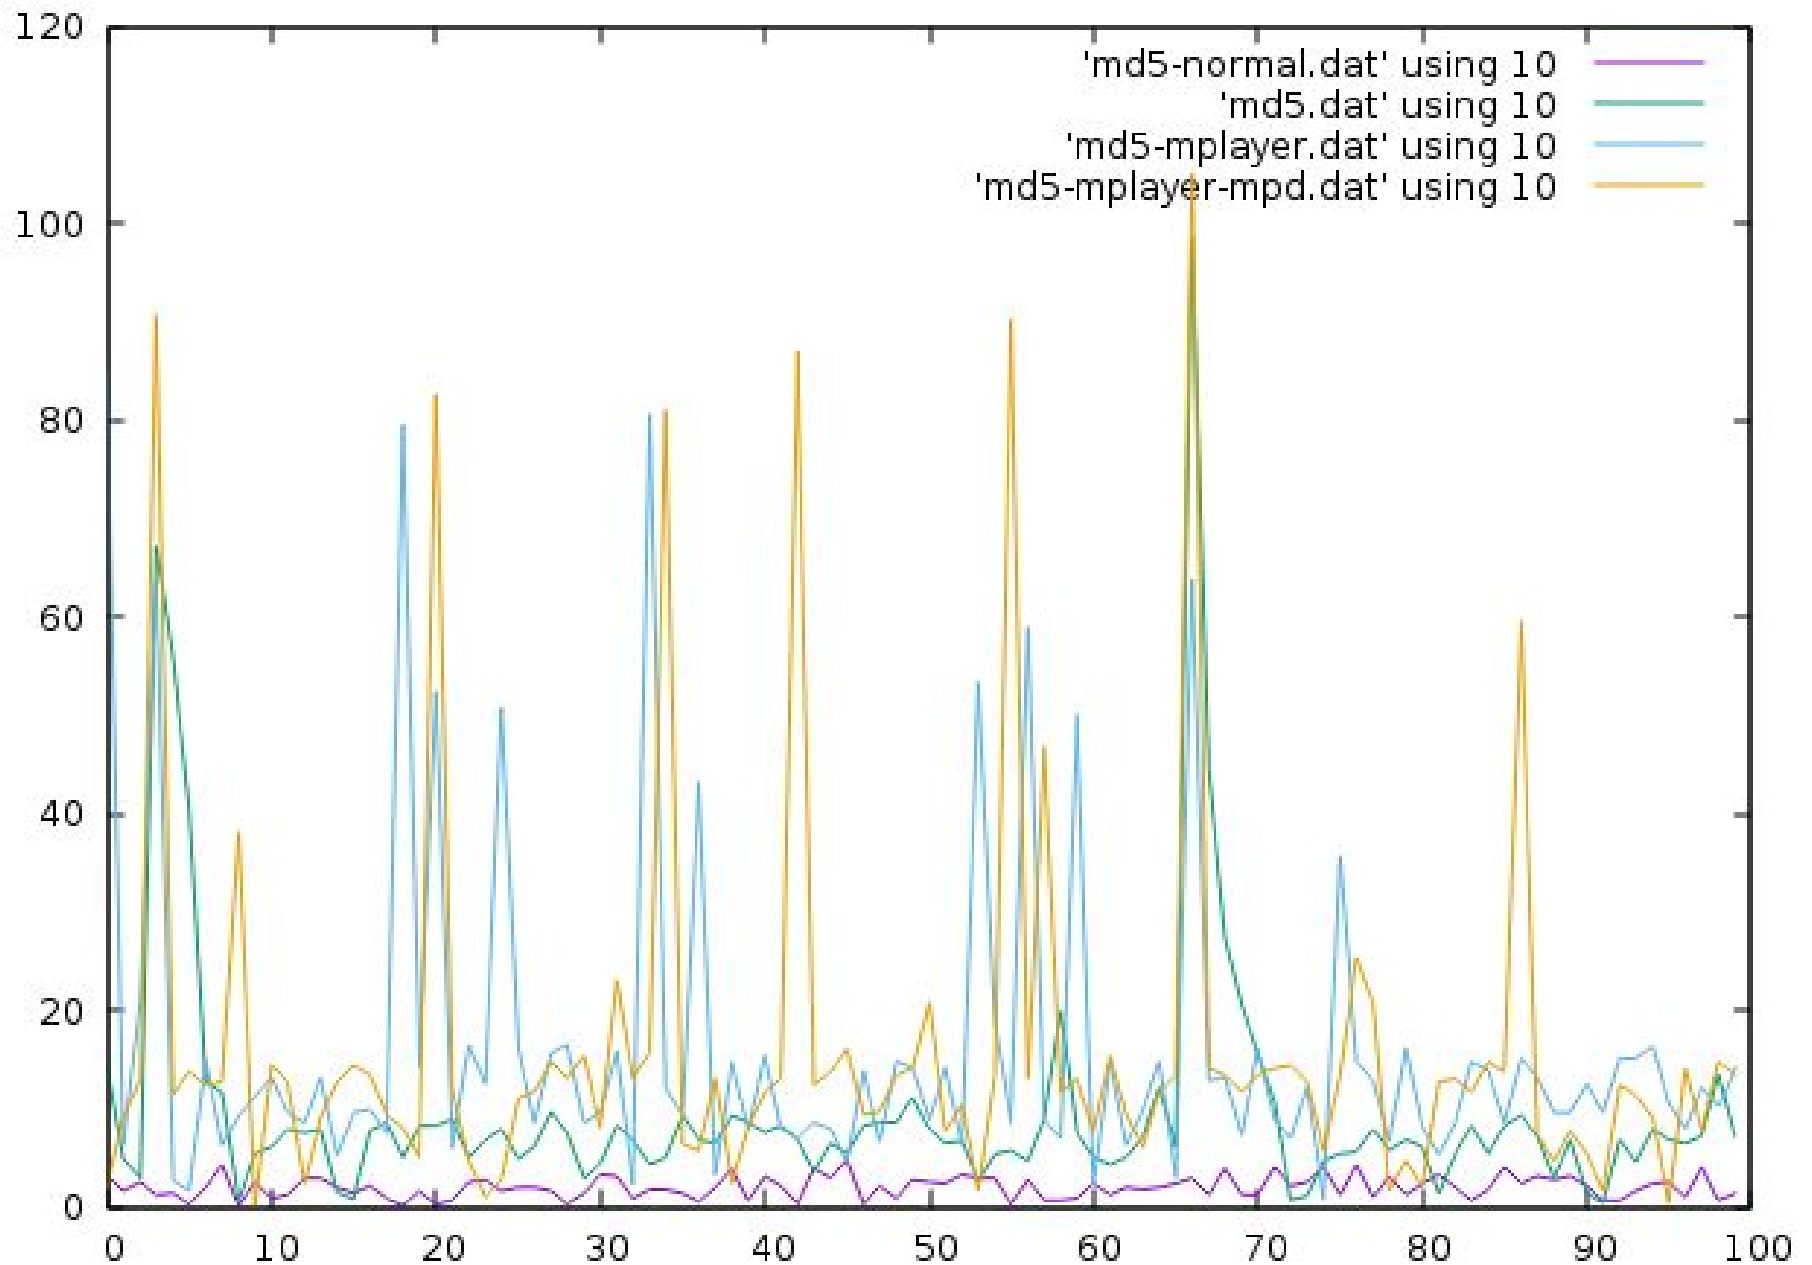


# Stability

- Accuracy of benchmarks depend on the ambient load on the machines resources.
- Most machines are not stable.
- Especially developer machines!



# MD5 hashing algorithm





Don't watch films and play music when  
benchmarking.

# Observer Effect



“In science, the term observer effect refers to changes that the act of observation will make on a phenomenon being observed.”

- Wikipedia

# Observer Effect

- Is not the Uncertainty Principle
- Logging
- Profiling
- Benchmarking



# Measuring time takes time.

```
1 <?php
2
3 for ($i = 0; $i < 10; $i++) {
4     $start = microtime(true);
5     $end = microtime(true);
6     echo $i . ': ' . ($end - $start) * (10 ** 6);
7     echo PHP_EOL;
8 }
9
```


```
~/Tmp >>> php test.php
0: 1.9073486328125
1: 0
2: 1.1920928955078
3: 0.95367431640625
4: 0.95367431640625
5: 1.1920928955078
6: 0
7: 0
8: 0.95367431640625
9: 0
```

# Recovery


- Sometimes external services may need time to calm down between iterations.
- But many don't.

# Sleep with Jackalope DBAL

- 50 iterations
- Querying 1000 nodes
- Sleep doesn't make much difference here.



# Sleep with Jackalope DBAL






# Conclusions

- We benchmark using microtime, because that's the best we have.
- Benchmarking can help you make smart decisions.
- A certain margin of error is unavoidable.
- Resource stability is important.

# Questions?

# PHPBench



“PHPBench is a Benchmark Runner for PHP  
which can Generate Reports”

- J. Lennon

# What is a Benchmark Runner?

Program that runs *a series of processes* which return the time (and memory) that it took to execute a specified piece of code.

# Writing benchmarks in PHPBench



# Benchmark Classes

- Must be suffixed with “Bench”.
- Are similar to PHPUnit
- Have their own autoloading env.
- Do not depend on the PHPBench library (not abstract classes, no interfaces).

# Benchmarking Hash Algorithms

```
1 <?php
2
3 class HashBench
4 {
5     public function benchMd5()
6     {
7         hash('md5', 'hello world');
8     }
9
10    public function benchSha256()
11    {
12        hash('sha256', 'hello world');
13    }
14 }
```

# Benchmarking Hash Algorithms

```
~/w/livecoding >>> phpbench run HashBench.php --report=default  
PhpBench 0.7.0-dev. Running benchmarks.
```

```
..  
Done (2 subjects, 2 iterations, 0 rejects) in 0.14s
```

benchmark	subject	group	params	revs	iter	rej	time	memory	deviation
HashBench	benchMd5		[]	1	0	0	11.0000µs	576b	0.00%
HashBench	benchSha256		[]	1	0	0	11.0000µs	576b	0.00%
					stability		100.00%		
					average	0.00	11.0000µs	576b	
					sum	0.00	22.0000µs	1,152b	

# Revolutions



# Revolutions

```
1 <?php
2
3 $revolutions = 10000;
4
5 $start = microtime(true);
6
7 for ($i = 0; $i < $revolutions; $i++) {
8     hash('md5', 'hello world');
9 }
10
11 $end = microtime(true);
12
13 $time = $end - $start;
```

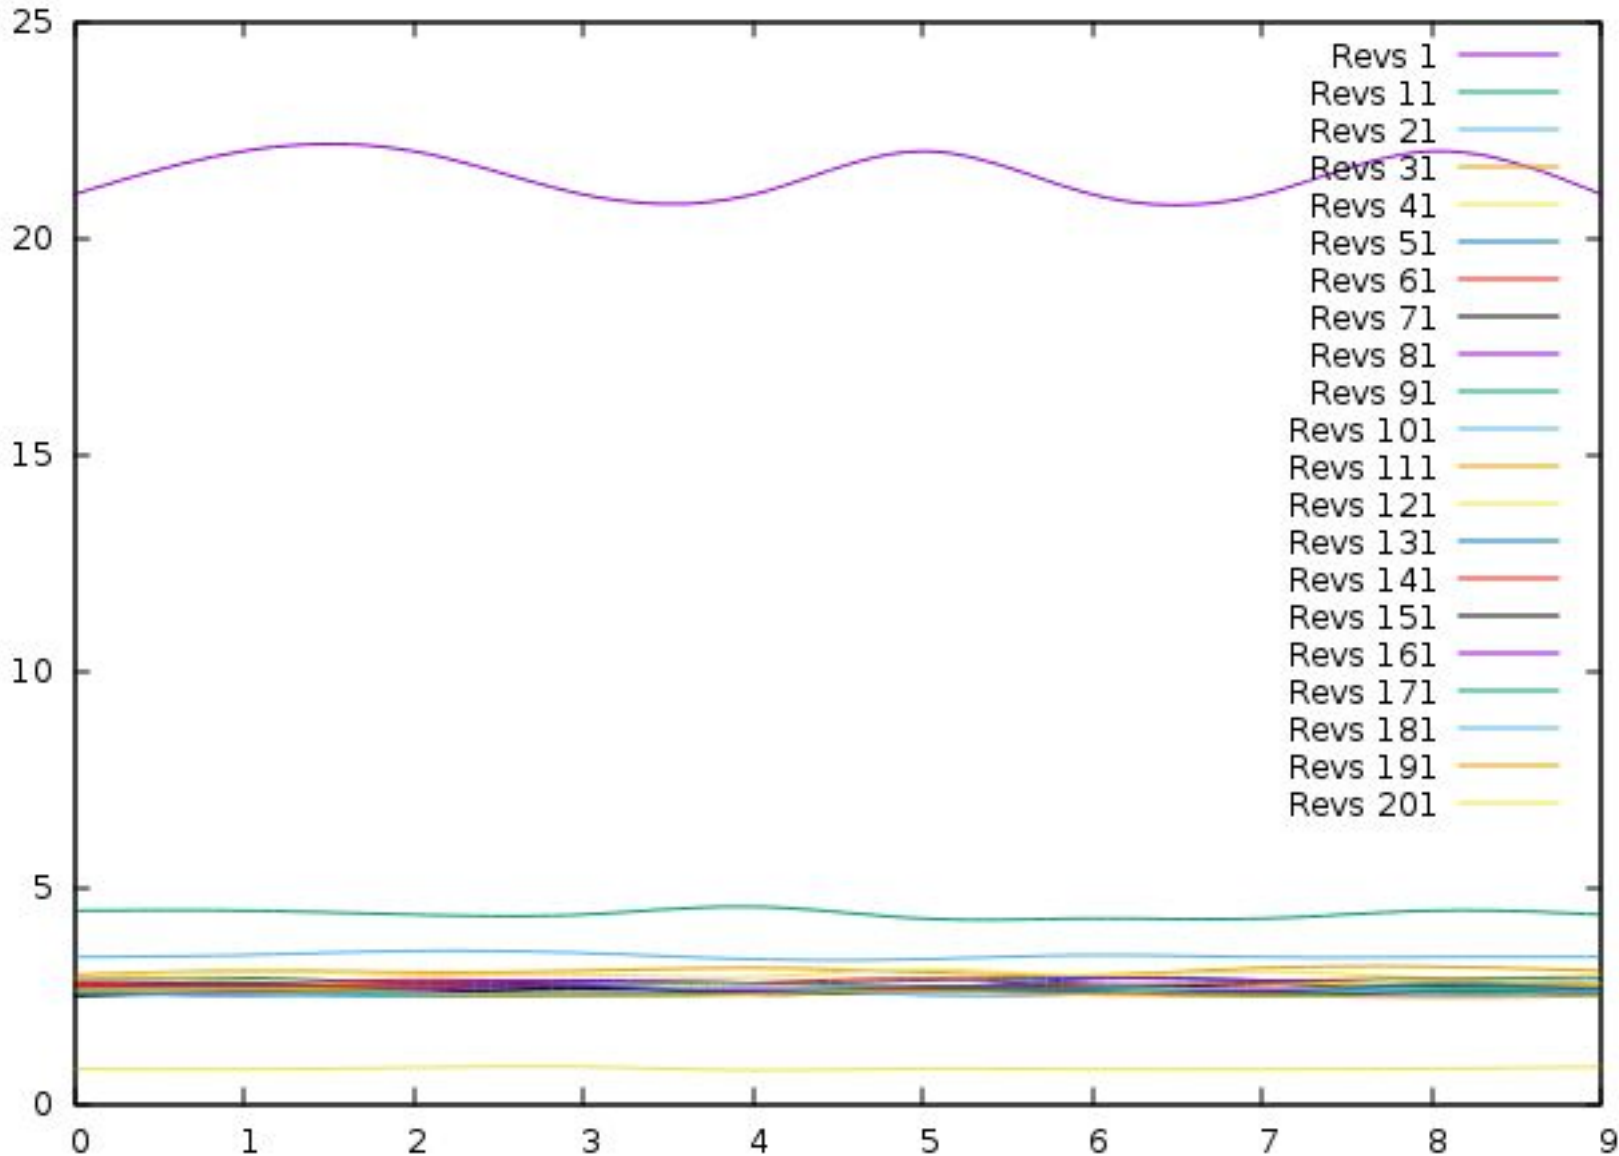


# Revolutions Increase Precision

$$time = total\ time / revolutions$$

Revolutions	MD5 Time	SHA256 Time
1	9ms	9ms
10	4.6ms	5.00ms
100	2.7ms	3.09ms
1000	1.87ms	2.20ms
10000	1.62ms	2.25ms

# Revolutions by iteration against time.



# Specifying Revolutions

```
1 <?php
2
3 /**
4  * @Revs(1000)
5  */
6 class HashBench
```

*Default revolutions for class*


```
8      /**
9       * @Revs(2000)
10      */
11     public function benchMd5()
12     {
13         hash('md5', 'hello world');
14     }
```

*Revolutions per method*

```
./bin/phpbench run examples/HashBench.php --report=default --revs=1000
```

*Revolutions overridden on command line*

# Iterations and Stability



# Iterations Confirm Stability

```
PhpBench 0.7.0-dev. Running benchmarks.
```

```
Done (1 subjects, 4 iterations, 0 rejects) in 0.38s
```

benchmark	subject	group	params	revs	iter	rej	time	memory	deviation
HashBench	benchMd5		[]	10000	0	0	3.0043µs	576b	3.09%
HashBench	benchMd5		[]	10000	1	0	3.0631µs	576b	1.19%
HashBench	benchMd5		[]	10000	2	0	3.1620µs	576b	2.00%
HashBench	benchMd5		[]	10000	3	0	3.1708µs	576b	2.28%
					stability		94.46%		
					average	0.00	3.1001µs	576b	
					sum	0.00	12.4002µs	2,304b	

## Stability 94.46%



# Iterations Confirm Stability

PhpBench 0.7.0-dev. Running benchmarks.

Done (1 subjects, 4 iterations, 0 rejects) in 0.37s

benchmark	subject	group	params	revs	iter	rej	time	memory	deviation
HashBench	benchMd5		[]	10000	0	0	3.2174µs	576b	12.17%
HashBench	benchMd5		[]	10000	1	0	3.2385µs	576b	12.91%
HashBench	benchMd5		[]	10000	2	0	1.7709µs	576b	38.26%
HashBench	benchMd5		[]	10000	3	0	3.2463µs	576b	13.18%
					stability		16.69%		
					average	0.00	2.8683µs	576b	
					sum	0.00	11.4731µs	2,304b	

*bad stability*

## Stability 16.69%

# Specifying Iterations

```
1 <?php
2
3 /**
4  * @Revs(10000)
5  * @Iterations(4)
6  */
7 class HashBench
8 {
9     public function benchMd5()
10    {
11        hash('md5', 'hello world');
12    }
13
14    public function benchSha256()
15    {
16        hash('sha256', 'hello world');
17    }
18 }
```

*as a class or method annotation*


```
./bin/phpbench run examples/HashBench.php --report=default --iterations=10
```

*as a command line override*


# Enforcing Stability: RT

- Specify a retry threshold (RT).
- Rerun the iteration set until all measurements fall within this tolerance level.
- Stability at the cost of time.
- Can be set either from configuration the command line.

# Retry count by retry threshold




# Effect of retry threshold feature





# Benchmarks Can Be Slow



# Benchmarking Slow Things

Indexing a large number of documents  
in different search implementations

```
~/w/m/MassiveSearchBundle >>> ./vendor/bin/phpbench run --config=phpbench.json --dump-file=report.xml
PhpBench 0.5. Running benchmarks.
Using configuration file: phpbench.json
.....
Done (6 subjects, 18 iterations) in 106.91s
Dumped result to report.xml
```

1 minute 46 seconds!

# Dumping the Results to XML

```
~/w/m/MassiveSearchBundle >>> ./vendor/bin/phpbench run --config=phpbench.json --dump-file=report.xml  
PhpBench 0.5. Running benchmarks.  
Using configuration file: phpbench.json  
  
.....  
Done (6 subjects, 18 iterations) in 106.91s  
Dumped result to report.xml
```

Dumped result to report.xml

# Reporting on the XML file


```
./vendor/bin/phpbench report report.xml --report=aggregate
```

```
~/w/m/MassiveSearchBundle >>> ./vendor/bin/phpbench report report.xml --report=aggregate
```

benchmark	subject	group	params	revs	iters	time	memory	deviation	stability
TestBench	benchIndex		[]	1	1	465,830.00µs	6,914,832b	0.00%	100.00%
ZendLuceneBench	benchIndex		[]	1	1	45,079,632.00µs	8,477,904b	+9,577.27%	100.00%
ElasticBench	benchIndex		[]	1	1	4,141,092.00µs	7,952,808b	+788.97%	100.00%

## Instant Report!

# Configuration



# JSON Configuration File

```
{  
    "bootstrap": "../vendor/autoload.php",  
    "path": "./",  
    "reports": {  
    }  
}
```

- Automatically uses phpbench.json if present
- Config can be specified using --config



# Reporting



# Reporting

- **Generators:** Produce a report document from a suite result document.
- **Reports:** Configurations for generators
- **Tabular:** Configuration language for creating reports.
- **Renderers:** For rendering to different output mediums.

# Generators

```
interface ReportGeneratorInterface
{
    /**
     * Generate the report.
     *
     * @param SuiteDocument $collection
     * @param array $config
     */
    public function generate(SuiteDocument $collection, array $config);
}
```

*(simplified)*

- **ConsoleTableGenerator**: Generates easy reports
- **ConsoleTableCustomGenerator**: Definition of complex reports
- **CompositeGenerator**: Generate multiple reports

# Reports

Reports are configured using an array

```
{
  "generator": "console_table_custom",
  "title": "Comparison of array location functions",
  "description": "This benchmark creates an array with
x of the current revolution. (or in the case of in_arrey,
  "file": "reports/array_keys.json"
},
```

*report using a specified generator*

```
{
  "extends": "aggregate",
  "title": "Cost of Setting",
  "description": "Comparison of different ways of setting properties",
  "selector": "//subject[group/@name='cost_of_setting']//variant",
  "exclude": ["benchmark"]
},
```

*report extending another report*

# Reports

- Can be named and defined in the PHPBench configuration
- Can be specified directly on the command line:

```
~/w/p/phpbench >>> ./bin/phpbench run examples/HashBench.php \
  --report='{ "extends": "aggregate", "exclude": ["benchmark", "subject", "group", "params"] }'
PhpBench 0.5. Running benchmarks.
..
Done (2 subjects, 20 iterations) in 0.54s
```

revs	iters	time	memory	deviation	stability
10000	10	2.65µs	616b	0.00%	86.54%
10000	10	3.10µs	592b	+17.16%	87.77%

# The Default Reports



# Default Report

```
~/w/p/phpbench >>> ./bin/phpbench run examples/HashBench.php --report=default --iterations=10 --revs=1000
PhpBench 0.5. Running benchmarks.
```

```
..
Done (2 subjects, 20 iterations) in 0.55s
```

benchmark	subject	group	params	revs	iter	time	memory	deviation
HashingBenchmark	benchMd5		[]	1000	0	2.55µs	576b	-8.34
HashingBenchmark	benchMd5		[]	1000	1	2.56µs	576b	-7.98
HashingBenchmark	benchMd5		[]	1000	2	2.66µs	576b	-4.3
HashingBenchmark	benchMd5		[]	1000	3	2.93µs	576b	+5.56%
HashingBenchmark	benchMd5		[]	1000	4	2.45µs	576b	-11.83
HashingBenchmark	benchMd5		[]	1000	5	2.42µs	576b	-12.83
HashingBenchmark	benchMd5		[]	1000	6	2.51µs	576b	-9.56
HashingBenchmark	benchMd5		[]	1000	7	2.65µs	576b	-4.63
HashingBenchmark	benchMd5		[]	1000	8	2.98µs	576b	+7.07%
HashingBenchmark	benchMd5		[]	1000	9	2.66µs	576b	-4.2
HashingBenchmark	benchSha256		[]	1000	0	2.96µs	576b	+6.35%
HashingBenchmark	benchSha256		[]	1000	1	3.04µs	576b	+9.48%
HashingBenchmark	benchSha256		[]	1000	2	2.80µs	576b	+0.81%
HashingBenchmark	benchSha256		[]	1000	3	2.84µs	576b	+2.21%
HashingBenchmark	benchSha256		[]	1000	4	2.77µs	576b	-0.35
HashingBenchmark	benchSha256		[]	1000	5	2.92µs	576b	+4.91%
HashingBenchmark	benchSha256		[]	1000	6	3.00µs	576b	+7.90%
HashingBenchmark	benchSha256		[]	1000	7	2.82µs	576b	+1.42%
HashingBenchmark	benchSha256		[]	1000	8	3.19µs	576b	+14.73%
HashingBenchmark	benchSha256		[]	1000	9	2.88µs	576b	+3.58%
stability >>						68.37%		
average >>						2.78µs	576b	
sum >>						55.57µs	11,520b	

# Aggregate Report

```
~/w/p/phpbench >>> ./bin/phpbench run examples/HashBench.php --report=aggregate --iterations=10 --revs=1000
PhpBench 0.5. Running benchmarks.
```

```
..
```

```
Done (2 subjects, 20 iterations) in 0.63s
```

benchmark	subject	group	params	revs	iters	time	memory	deviation	stability
HashingBenchmark	benchMd5		[]	10000	10	2.90µs	576b	0.00%	85.26%
HashingBenchmark	benchSha256		[]	10000	10	3.71µs	576b	+27.87%	70.04%

# Plain Report

benchmark	subject	group	params	revs	iter	rej	time	memory	deviation
SortingBench	benchSort	sort	[]	1000	0	0	198.412	584	0.5636087176888
SortingBench	benchSort	sort	[]	1000	1	0	198.272	584	0.49265078560568
SortingBench	benchSort	sort	[]	1000	2	0	196.188	584	0.5636087176888
SortingBench	benchSort	sort	[]	1000	3	0	196.328	584	0.49265078560568
SortingBench	benchNatSort	sort	[]	1000	0	0	184.857	584	0.013119969269515
SortingBench	benchNatSort	sort	[]	1000	1	0	184.429	584	0.21844072546667
SortingBench	benchNatSort	sort	[]	1000	2	0	185.031	584	0.10725913021367
SortingBench	benchNatSort	sort	[]	1000	3	0	185.014	584	0.09806162598349
SortingBench	benchASort	sort	[]	1000	0	0	192.099	584	0.6251697425087
SortingBench	benchASort	sort	[]	1000	1	0	191.514	584	0.92779638658614
SortingBench	benchASort	sort	[]	1000	2	0	197.536	584	2.1874474606521
SortingBench	benchASort	sort	[]	1000	3	0	192.081	584	0.63448133155723

using console renderer

SortingBench	benchSort	sort	[]	1000	0	0	197.827	584	6.8291987957269
SortingBench	benchSort	sort	[]	1000	1	0	255.247	584	20.213962173956
SortingBench	benchSort	sort	[]	1000	2	0	199.439	584	6.0699933710817
SortingBench	benchSort	sort	[]	1000	3	0	196.796	584	7.314770007147
SortingBench	benchNatSort	sort	[]	1000	0	0	188.78	584	1.4661204043488
SortingBench	benchNatSort	sort	[]	1000	1	0	185.058	584	0.53439289231923
SortingBench	benchNatSort	sort	[]	1000	2	0	185.243	584	0.43495845925002
SortingBench	benchNatSort	sort	[]	1000	3	0	185.128	584	0.49676905277953
SortingBench	benchASort	sort	[]	1000	0	0	193.645	584	0.9078817265516
SortingBench	benchASort	sort	[]	1000	1	0	190.724	584	0.61424341235339
SortingBench	benchASort	sort	[]	1000	2	0	191.096	584	0.42039522622787
SortingBench	benchASort	sort	[]	1000	3	0	192.146	584	0.12675691202966

using delimited renderer

# Custom Reports

```
1 {
2   "includes": [
3     [ "console_classes.json", ["classes"] ]
4   ],
5   "rows": [
6     {
7       "group": "body",
8       "cells": [
9         {
10          "name": "time",
11          "class": "time",
12          "expr": "number(descendant-or-self::iteration/@time) div"
13        },
14        {
15          "name": "memory",
16          "class": "memory",
17          "expr": "number(descendant-or-self::iteration/@memory)"
18        },
19        {
20          "name": "deviation",
21          "class": "deviation",
22          "expr": "deviation(average(//iteration/@time), number(de"
23        )
24      ],
25      "with_query": "//iteration"
26    },
27    {
28      "cells": [
29
```

- New reports can be registered in the configuration file.
- The “console\_table\_custom” generator allows you to specify a custom Tabular report definition.

# Tabular



The image displays a collage of overlapping spreadsheets, each containing numerical data. The spreadsheets are arranged in a way that they appear to be floating and overlapping each other, creating a sense of depth. The data is presented in a tabular format, with rows and columns of numbers. The numbers are in various colors, including blue, black, and red, and are of different sizes. The overall effect is a dense, layered representation of data.

451	368
164	94
166	172

46	83	74	29	10
73	38	99	25	73
91	85	40	78	49
30	62	49	32	31

340	301	336	293	317
232	377	431	411	451
430	451	367	439	184
182	139	144	235	186

433	896	2.132
.870	2.845	1.001
2.427	1.133	1.308
2.424	2.697	1.710
1.692	1.844	1.725
1.199	1.903	1.442
2.032	1.198	2.453

2.390	3.850	2.175	1.389	2.833	3.928
1.920	1.748	2.387	2.930	1.389	1.253
3.928	3.176	2.514	2.635	2.119	1.373
1.287	1.272	2.303	2.738	2.115	2.001
2.110	1.928	1.902	1.627	2.738	2.353
3.292	3.393	2.990	2.117	2.517	2.001
1.272	1.928	1.897	2.119	2.519	3.555

110	383	212	259	555
839	154	755	559	555
748	825	339	555	555


290	92	266
243	430	159
249	277	324
175	304	

809	2.402
1.988	

# Tabular

- Developed for PHPBench but usable anywhere.
- Library for creating tabular *data*
- Accepts any XML document as its input
- Uses Xpath queries and expressions.
- Transforms it according to a Tabular JSON given definition.





# Given an XML file

```
<?xml version="1.0"?>
<store>
  <book>
    <title>War and Peace</title>
    <price>5.00</price>
  </book>
  <book>
    <title>One Hundered Years of Soliture</title>
    <price>7</price>
  </book>
</store>
```

# And a Tabular definition

```
{
  "rows": [
    {
      "cells": [
        {
          "name": "title",
          "expr": "string(/title)"
        },
        {
          "name": "price",
          "expr": "number(/price)"
        }
      ],
      "with_query": "//book"
    },
    {
      "cells": [
        {
          "name": "price",
          "expr": "sum(/price)"
        }
      ]
    }
  ]
}
```

# It generates data

Book	Price
War and Peace	5
One Hundred Years of Solitude	7
	12

- Generated data returned as an XML document
- Does not render tables

# It does other stuff!

- Formats cells
- Iterate over Xpath queries
- Iterate over items
- “Compiler” passes
- Split tables into distinct groups (e.g. header, body, footer).
- Sorting
- Include other definitions

# Renderers

```
16 /**
17  * Implementors render the DOM Document generated by a generator to some visual medium.
18  *
19  * This might be either direct output (e.g. to the console, streaming markup) or to a file.
20  *
21  * Example implementors might be XsltRenderer, ConsoleRenderer, etc
22  */
23 interface RendererInterface extends ConfigurableInterface
24 {
25     public function render(Document $report, array $config);
26
27     public function getDefaultOutputs();
28 }
```

- Render reports to different mediums
- Console
- HTML
- Markdown
- Delimited file (csv, tsv, etc)

# HTML Output

## PHPBench Benchmark Results

benchmark	subject	group	params	revs	its	rej	time	memory	deviation	stability
HashingBenchmark	benchMd5	hashing	[]	10000	10	0	2.6861µs	576b	0.00%	79.45%
HashingBenchmark	benchSha256	hashing	[]	10000	10	0	3.0700µs	576b	14.29%	80.72%
HashingBenchmark	benchSha1	hashing	[]	10000	10	0	2.9019µs	576b	8.03%	77.14%

Generated 2015-11-19 12:07:17 by [PHPBench](#) v0.6



- XSLT
- Specify custom layout



# Markdown Output

The screenshot shows a GitHub repository page for 'phpbench / phpbench'. The repository has 11 Unwatch, 222 Stars, and 6 Forks. The 'Issues' tab is selected, showing 15 issues. A specific issue, 'Example #207', is highlighted, opened by 'dantleech' just now, with 0 comments. The issue content displays a table titled 'PHPBench Benchmark Results' with the following data:

benchmark	subject	group	params	revs	its	rej	time
HashingBenchmark	benchMd5	hashing	[]	10000	10	0	2.6799μs
HashingBenchmark	benchSha256	hashing	[]	10000	10	0	3.0908μs
HashingBenchmark	benchSha1	hashing	[]	10000	10	0	2.9449μs

On the right side of the issue, there are sections for 'Labels' (None yet), 'Milestone' (No milestone), 'Assignee' (No one—assign yourself), and 'Notifications' (Unsubscribe button). A note at the bottom states: 'You're receiving notifications because you authored the thread.'

- Github flavored

# Delimited Output

- Output to a delimited list
- To file or STDOUT
- Import into:
  - GNUPlot (make graphs)
  - Spreadsheets
  - Whatever ..

# Conclusion

- Benchmarking can help you make ***smart decisions***.
- PHPBench can be used to obtain more ***reliable*** benchmark results.
- Tabular can be used to generate ***sexy*** reports.

# Documentation

PHPBench 0.5 documentation »

Table Of Contents

Welcome to PHPBench's documentation!  
Indices and tables

Next topic

Introduction

This Page

Show Source

Quick search

Go

Enter search terms or a module, class or function name.

## Welcome to PHPBench's documentation!

PHPBench is a benchmarking framework for PHP. Find it on [Github](#).

Contents:

- Introduction
  - Why PHPBench?
  - Are There Other Benchmarking Frameworks?
- Installing
  - Composer Install
  - Composer Global Install
- Quick Start
  - Create your project
  - PHPBench configuration
  - Creating and running a benchmark
  - Increase Stability
  - Customize Reports
  - Configuration
- Writing Benchmarks
  - Improving Precision: Revolutions
  - Verifying and Improving Stability: Iterations
  - Establishing State: Before and After
  - Parameterized Benchmarks
  - Groups

[phpbench.readthedocs.org](http://phpbench.readthedocs.org)

# Questions?

# Live Coding!

# The End

- Twitter: **@phpbench @dantleech**
- Website: [www.dantleech.com](http://www.dantleech.com)
- PHPBench: [github.com/phpbench/phpbench](https://github.com/phpbench/phpbench)
- Docs: [phpbench.readthedocs.org](http://phpbench.readthedocs.org)
- Tabular: [github.com/phpbench/tabular](https://github.com/phpbench/tabular)

Suggestions welcome – PHPBench and this talk are works in progress :)

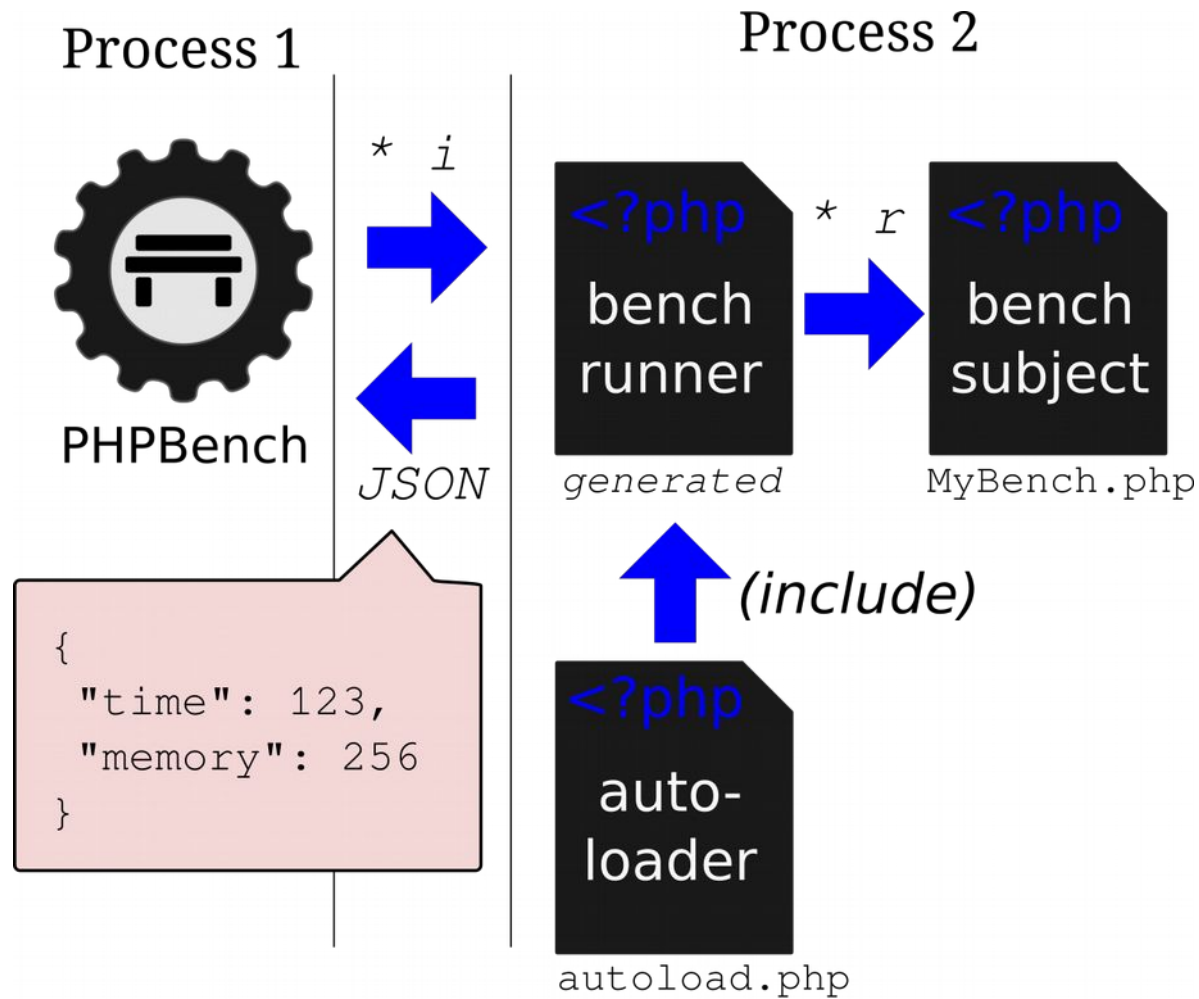


Extra

# Benchmark Process Isolation

- Benchmarks executed by a generated script in a separate process.
- PHPBench does add additional overhead to your benchmarks.

# How it works



*i* = iterations, *r* = revolutions

# The Benchmarking Script

- Generated in /tmp
- Executed with the Symfony Process component

# PHPBench Benchmarking Script

```
1 <?php
2
3 gc_disable();
4
5 $class = '{{ class }}';
6 $file = '{{ file }}';
7 $subject = '{{ subject }}';
8 $revolutions = {{ revolutions }};
9 $bootstrap = '{{ bootstrap }}';
10
11 if ($bootstrap) {
12     require_once($bootstrap);
13 }
14
15 require_once($file);
16
17 $benchmark = new $class();
18 $startMemory = memory_get_usage();
19 $startTime = microtime(true);
20
21 for ($i = 0; $i < $revolutions; $i++) {
22     $benchmark->$subject();
23 }
24
25 $endTime = microtime(true);
26 $endMemory = memory_get_usage();
27
28 echo json_encode(array(
29     'memory' => $endMemory - $startMemory,
30     'time' => ($endTime * 1000000) - ($startTime * 1000000),
31 ));
32
```

*script template (simplified)*

- Does not require PHPBench to be autoloaded
- Generated in tmp directory
- Disables Garbage Collection

# What About Profiling?

- Provides detailed analysis of the whole lifecycle, including timings
- Essential for a deeper understanding of code performance
- Tools
  - Xdebug with a visualiser (e.g. KCacheGrind)
  - Blackfire.io

# Benchmarking vs. Profiling

- Profiling (using tools such as Xdebug and Blackfire) measures time taken per function call.
- It tells you which parts of your code are slow.
- Benchmarking measures time taken to do something specific.
- It helps you to compare different ways of doing something.



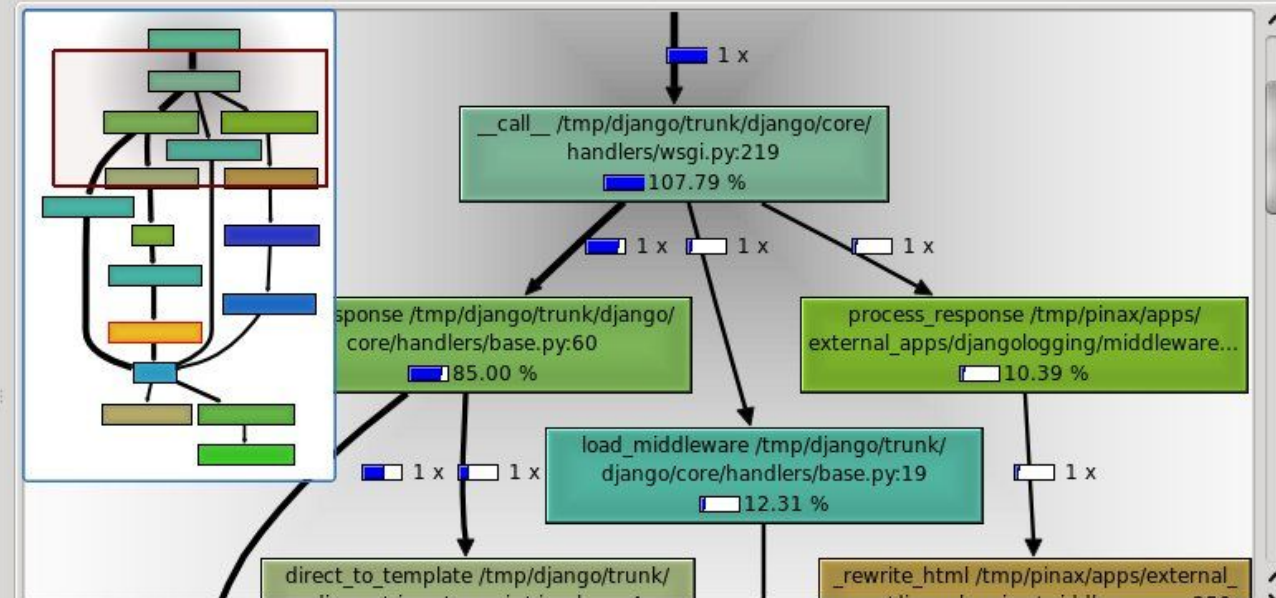
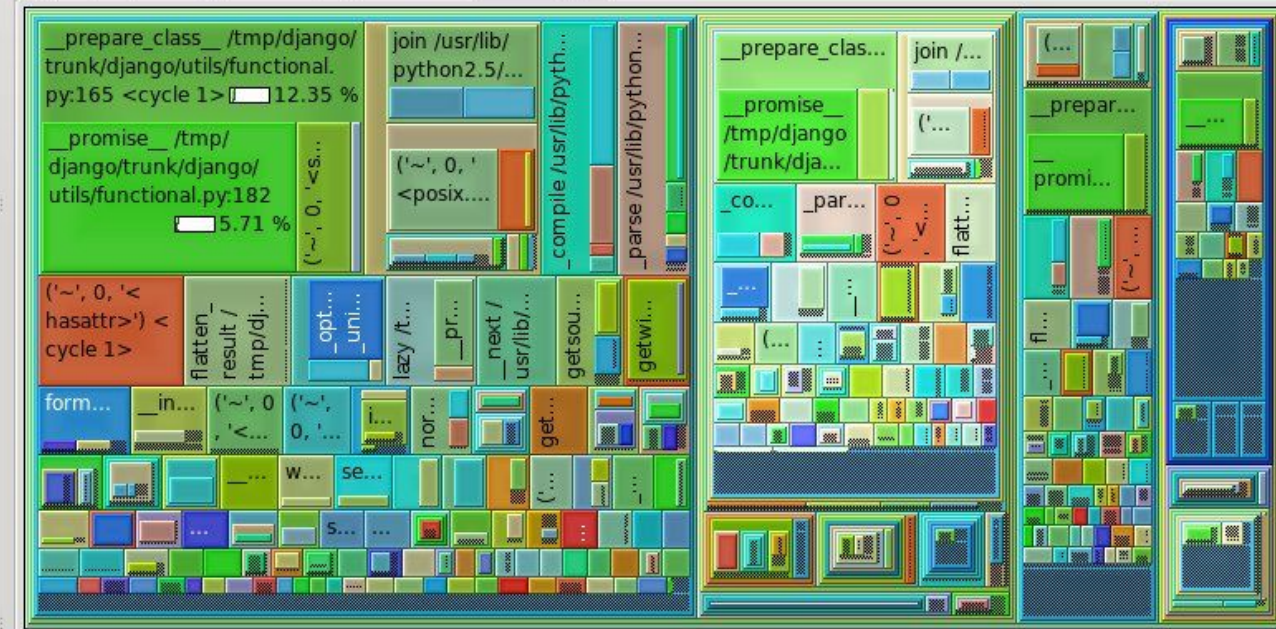
# Flat Profile

Search: (No Grouping)

Incl.	Self	Called	Function	Location
107.79	0.00	1	<code>_call_ /tmp/django/trunk/d... wsgi.py</code>	
107.79	0.00	(0)	<code>_call_ /tmp/django/trunk/d... basehttp.py</code>	
96.63	54.79	4 843	<code>&lt;cycle 1&gt;</code>	(unknown)
85.00	0.00	1	<code>get_response /tmp/django/tr... base.py</code>	
57.44	0.00	1	<code>resolve /tmp/django/trunk/dj... urlresolvers.</code>	
27.85	0.19	3	<code>&lt;cycle 4&gt;</code>	(unknown)
27.39	0.00	1	<code>direct_to_template /tmp/dja... simple.py</code>	
24.08	0.02	243	<code>render_node /tmp/django/tr... debug.py</code>	
23.77	0.00	8	<code>render /tmp/django/trunk/dj... defaulttags.</code>	
20.78	8.72	675	<code>_prepare_class_ /tmp/djan... functional.p</code>	
12.31	0.00	1	<code>load_middleware /tmp/djang... base.py</code>	
11.29	1.82	868	<code>realpath /usr/lib/python2.5/... posixpath.p</code>	
10.39	0.00	1	<code>process_response /tmp/pina... middleware</code>	
10.37	0.00	1	<code>rewrite_html /tmp/pinax/ap... middleware</code>	
9.61	9.61	36 450	<code>_promise_ /tmp/django/tru... functional.p</code>	
7.67	0.00	5	<code>sql_to_html /tmp/pinax/apps... middleware</code>	
7.67	0.00	1	<code>get_and_clear_records /tm... middleware</code>	
4.88	3.35	1 907	<code>_compile /usr/lib/python2.5/... sre_compile</code>	
4.66	3.33	1 449	<code>_parse /usr/lib/python2.5/sr... sre_parse.p</code>	
4.22	1.24	5 216	<code>islink /usr/lib/python2.5/posi... posixpath.p</code>	
4.05	4.05	45 720	<code>('~', 0, '&lt;hasattr&gt;') &lt;cycle ... ~</code>	
3.88	2.65	5 577	<code>join /usr/lib/python2.5/posix... posixpath.p</code>	
2.99	2.91	216	<code>flatten_result /tmp/django/tr... regex_help</code>	
2.62	0.68	357	<code>_optimize_charset /usr/lib/py... sre_compile</code>	
2.57	0.00	2	<code>render_to_string /tmp/djang... loader.py</code>	
2.55	0.00	5	<code>_init_ /tmp/django/trunk/d... _init_.py</code>	
2.55	0.00	5	<code>compile_string /tmp/django/... _init_.py</code>	
2.52	1.46	675	<code>lazy /tmp/django/trunk/djan... functional.p</code>	
2.52	0.00	5	<code>get_template_from_string /t... loader.py</code>	
2.23	2.23	5 216	<code>('~', 0, '&lt;posix.lstat&gt;') ~</code>	
2.21	2.21	15 798	<code>_next /usr/lib/python2.5/sr... sre_parse.p</code>	
2.14	2.16	38 757	<code>('~', 0, '&lt;setattr&gt;') ~</code>	
1.89	1.19	602	<code>getsourcefile /usr/lib/python... inspect.py</code>	
1.85	0.27	331	<code>_compile_info /usr/lib/pytho... sre_compile</code>	
1.85	0.27	5	<code>&lt;cycle 3&gt;</code>	(unknown)
1.82	0.15	1 458	<code>abspath /usr/lib/python2.5/... posixpath.p</code>	
1.70	0.27	64	<code>parse /tmp/django/trunk/dja... _init_.py</code>	

## \_call\_ /tmp/django/trunk/django/core/handlers/wsgi.py:219

Types Callers All Callers Source Code Callee Map



# Profiling vs. Benchmarking

- Profiling
  - Is necessarily slower.
  - Measurements are the result of a single execution
  - Feedback is not instantaneous
- Benchmarking
  - Code runs at its natural speed
  - Units and Services can be timed in isolation
  - Feedback available instantaneously

# Continuous Benchmarking?

- Automated performance regression testing
- Would require a VM that runs at a constant speed.
- Travis CI does not currently meet this requirement.

# Calculating Stability

#	Time
1	1
2	1
3	1

= 100% stability



#	Time
1	1
2	0.5
3	1

= 50% stability



# Improving Stability

- Multiple iterations show correlations in samples.
- We can remove the sample with a deviation  $>$  a given threshold.
- We can then use the average of the remaining samples.

